

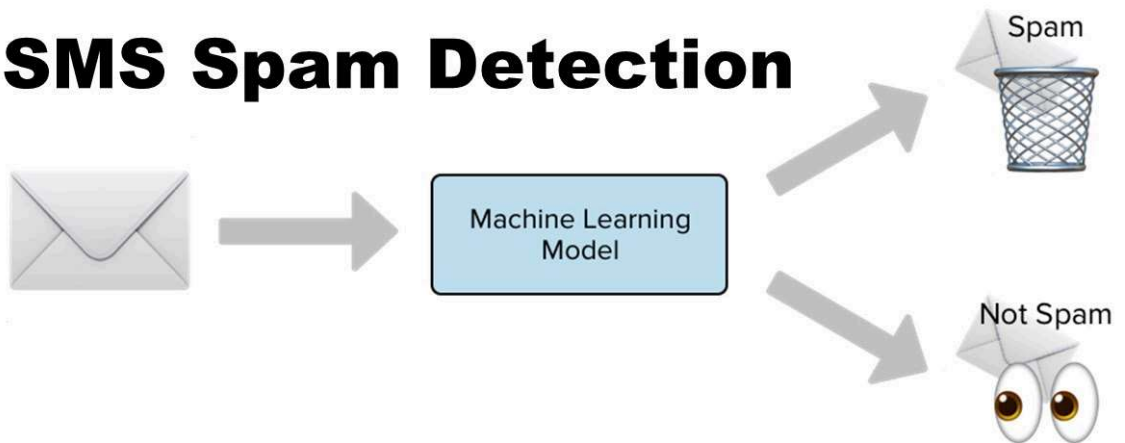
Spam_SMS_email_Classifier

-Mayank Srivastava



- (<https://linkedin.com/in/mayank-srivastava-6a8421105/>)

SMS Spam Detection



Dataset: <https://www.kaggle.com/datasets/uciml/sms-spam-collection-dataset> (<https://www.kaggle.com/datasets/uciml/sms-spam-collection-dataset>)

```
# 1. Data cleaning
# 2. EDA
# 3. Text Preprocessing
# 4. Model building
# 5. Evaluation
# 6. Improvement
# 7. Website
# 8. Deploy
```

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
```

```
In [2]: 1 df= pd.read_csv('spam.csv', encoding='latin-1')
        2 df.head()
```

Out[2]:

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

```
In [3]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0    v1          5572 non-null   object
1    v2          5572 non-null   object
2    Unnamed: 2   50 non-null     object
3    Unnamed: 3   12 non-null     object
4    Unnamed: 4   6 non-null      object
dtypes: object(5)
memory usage: 217.8+ KB

Observations:

- No Null values
- size = 5572 x 4
```

1. Data Cleaning

```
In [4]: 1 # Dropping the extra columns
2 df.drop(columns=['Unnamed: 2','Unnamed: 3','Unnamed: 4'], inplace = True)
3
4 #renaming the columns
5 df.rename(columns={'v1': 'label', 'v2': 'text'}, inplace = True)
6
7 # replacing ham -> 0 , spam -> 1
8 df.label.replace(['ham','spam'],[0,1], inplace = True)
9
10 df.head()
```

Out[4]:

	label	text
0	0	Go until jurong point, crazy.. Available only ...
1	0	Ok lar... Joking wif u oni...
2	1	Free entry in 2 a wkly comp to win FA Cup fina...
3	0	U dun say so early hor... U c already then say...
4	0	Nah I don't think he goes to usf, he lives aro...

```
In [5]: 1 # checking for duplicates
2
3 df.duplicated().sum()
```

Out[5]: 403

```
In [6]: 1 # removing duplicates
2 df =df[~df.duplicated()]
3 df.shape
```

Out[6]: (5169, 2)


```
In [7]: 1 5572-403
```

Out[7]: 5169

2. EDA

```
In [8]: 1 data= df.label.value_counts()
2 plt.pie(data.values, labels=['ham', 'spam'], autopct = "%1.2f%%", explode = [0.0,0.1])
3 plt.title('Data Distribution Spam (1) and Not-Spam (0)')
4 plt.show()
5 data
```

Data Distribution Spam (1) and Not-Spam (0)



```
Out[8]: label
0      4516
1       653
Name: count, dtype: int64
```

Feature Engineering

```
In [9]: 1 # Lets find no. of characters, words and sentences in spam and ham
2 import nltk
3 from nltk import word_tokenize,sent_tokenize
4
5 df['n_chars'] = df.text.str.len()
6 df['n_words'] =df.text.apply(lambda x: len(word_tokenize(x)))
7 df['n_sent'] =df.text.apply(lambda x: len(sent_tokenize(x)))
8 df.head()
```

```
Out[9]:
```

	label	text	n_chars	n_words	n_sent
0	0	Go until jurong point, crazy.. Available only ...	111	24	2
1	0	Ok lar... Joking wif u oni...	29	8	2
2	1	Free entry in 2 a wkly comp to win FA Cup fina...	155	37	2
3	0	U dun say so early hor... U c already then say...	49	13	1
4	0	Nah I don't think he goes to usf, he lives aro...	61	15	1

```
In [10]: 1 # plotting spam vs ham based on no. of characters
2 plt.figure(figsize =(10,4))
3 sns.histplot(df[df.label == 0]['n_chars'])
4 sns.histplot(df[df.label == 1]['n_chars'], color= 'red')
```


```
Out[10]: <Axes: xlabel='n_chars', ylabel='Count'>
```



Observation: Spam text has generally higher character count than ham

```
In [11]: 1 # plotting spam vs ham based on no. of words
2 plt.figure(figsize=(10,4))
3 sns.histplot(df[df.label == 0]['n_words'])
4 sns.histplot(df[df.label == 1]['n_words'], color='red')
```

Out[11]: <Axes: xlabel='n_words', ylabel='Count'>



Observation: Spam text has generally higher word count than ham

```
In [12]: 1 df.groupby('label')[['n_chars', 'n_words', 'n_sent']].mean().transpose().rename(columns={0:'ham', 1:'spam'})
```

Out[12]:

label	ham	spam
n_chars	70.459256	137.891271
n_words	17.123782	27.667688
n_sent	1.820195	2.970904

Observation- On an average:

- ham text has avg 70 characters while spam has almost double characters.
- ham text has 17 words, spam has 27 words
- ham has ~ 2 sentences, spam has ~3 sentences.

3. Text preprocessing

```
In [13]: 1 import nltk
2
3 nltk.download('stopwords')
4 nltk.download('punkt')
5 from nltk.corpus import stopwords
6
7 from nltk.stem import PorterStemmer
8 from nltk.stem import WordNetLemmatizer

[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\hp\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to
[nltk_data] C:\Users\hp\AppData\Roaming\nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

```
In [14]: 1 # Lets define a helper fuction for standardization of text
```

```
In [15]: 1 # remove punctuations
2 import string
3 punctuation= string.punctuation
4 punctuation
```

Out[15]: '!"#\$%&'()*+,-./:;<=>@[\\]^_`{|}~'

```
In [16]: 1 # stop words
2 from nltk.corpus import stopwords
3 stops = set(stopwords.words('english'))
4 stops
```

```
Out[16]: {'a',
'about',
'above',
'after',
'again',
'against',
'ain',
'all',
'am',
'an',
'and',
'any',
'are',
'aren',
"aren't",
'as',
'at',
'be',
'because',
'...
```

```
In [17]: 1 import re
2 stem = PorterStemmer()
3 lemm = WordNetLemmatizer()
4 def standardize(text):
5
6     # change to lower case
7     text = text.lower()
8
9     # keep only alpha numerics
10    # assuming _ sign is used for space in text, replacing it with space
11    text = re.sub('_', ' ', text)
12    text = re.findall(r"\w+", text)
13
14    # now text has been converted into list of words , after re.findall
15
16    # remove punctuations and stopwords
17    text = [i for i in text if i not in stops and i not in punctuation]
18
19    # Lemmatization
20    text = [lemm.lemmatize(i) for i in text]
21
22    # stemming
23    text = [stem.stem(i) for i in text]
24
25    return(' '.join(text))
```

```
In [18]: 1 standardize("Hello' there I am peter, from22nd_street the mayor's office, who__ are. you? Dear %. Dancing ate remote
```

```
Out[18]: 'hello peter from22ndstreet mayor offic dear danc ate remot'
```

```
In [19]: 1 df['std_text'] = df.text.apply(standardize)
2 df.std_text
```

```
Out[19]: 0      go jurong point crazi avail bugi n great world...
1              ok lar joke wif u oni
2      free entri 2 wkli comp win fa cup final tkt 21...
3              u dun say earli hor u c already say
4              nah think go usf life around though
        ...
5567    2nd time tri 2 contact u u â 750 pound prize 2...
5568              i b go esplanad fr home
5569              piti mood suggest
5570    guy bitch act like interest buy someth els nex...
5571              rofl true name
Name: std_text, Length: 5169, dtype: object
```


```
In [20]: 1 # pip install wordcloud
```

```
In [21]: 1 # word_cloud of spam messages
2 from wordcloud import WordCloud
3 wc = WordCloud(width =720, height =720, min_font_size =10, background_color = 'white')
```

```
In [22]: 1 spam_wc = wc.generate(df[df.label == 1]['std_text'].astype('str').str.cat(sep=" "))
2 # concatentaing texts with seperator space
```

```
In [23]: 1 plt.figure(figsize =(8,6))
          2 plt.imshow(spam wc)
```

```
Out[23]: <matplotlib.image.AxesImage at 0x232ef89a810>
```



```
In [24]: 1 ham_wc = wc.generate(df[df.label == 0]['std_text'].astype('str').str.cat(sep=" "))
          2 # concatentaing texts with separator space
```

```
In [25]: 1 plt.figure(figsize =(8,6))
          2 plt.imshow(ham wc)
```

```
Out[25]: <matplotlib.image.AxesImage at 0x232ef9ba490>
```



```
In [26]: 1 df[df.label == 1]['std_text']
```

```
Out[26]: 2      free entri 2 wkli comp win fa cup final tkt 21...
        5      freemsg hey darl 3 week word back like fun sti...
        8      winner valu network custom select receivea à 9...
        9      mobil 11 month u r entitil updat latest colour ...
        11     six chanc win cash 100 20 000 pound txt csh11 ...
           ...
        5537    want explicit sex 30 sec ring 02073162414 cost...
        5540    ask 3mobil 0870 chatlin inclu free min india c...
        5547    contract mobil 11 mnth latest motorola nokia e...
        5566    remind o2 get 2 50 pound free call credit deta...
        5567    2nd time tri 2 contact u u à 750 pound prize 2...
Name: std_text, Length: 653, dtype: object
```

```
In [27]: 1 # Top 30 words in spam,
2 # made a single list, used value_counts
3 spam_words = []
4 for i in df[df.label == 1]['std_text'].str.split(' '):
5     spam_words+=i
6 print(spam_words)

'1', 'minmobsmorelcpobox177hp51fl', 'urgent', 'tri', 'contact', 'u', 'today', 'draw', 'show', 'ã', '800', 'prize', 'g
uarante', 'call', '09050001295', 'land', 'line', 'claim', 'a21', 'valid', '12hr', 'monthli', 'password', 'wap', 'mobs
i', 'com', '391784', 'use', 'wap', 'phone', 'pc', 'today', 'vodafone', 'number', 'end', '0089', 'last', 'four', 'digi
t', 'select', 'receiv', 'ã', '350', 'award', 'number', 'match', 'pleas', 'call', '09063442151', 'claim', 'ã', '350',
'award', 'free', 'top', 'rington', 'sub', 'weekli', 'rington', 'get', '1st', 'week', 'free', 'send', 'subpoli', '8161
8', '3', 'per', 'week', 'stop', 'sm', '08718727870', 'free', 'msg', 'sorri', 'servic', 'order', '81303', 'could', 'de
liv', 'suffici', 'credit', 'pleas', 'top', 'receiv', 'servic', 'hard', 'live', '121', 'chat', '60p', 'min', 'choos',
'girl', 'connect', 'live', 'call', '09094646899', 'cheap', 'chat', 'uk', 'biggest', 'live', 'servic', 'vu', 'bcm1896w
c1n3xx', 'wow', 'boy', 'r', 'back', 'take', '2007', 'uk', 'tour', 'win', 'vip', 'ticket', 'pre', 'book', 'vip', 'clu
b', 'txt', 'club', '81303', 'trackmarqu', 'ltd', 'info', 'vipclub4u', 'hi', 'mandi', 'sullivan', 'call', 'hotmix', 'f
m', 'chosen', 'receiv', 'ã', '5000', '00', 'easter', 'prize', 'draw', 'pleas', 'telephon', '09041940223', 'claim', '2
9', '03', '05', 'prize', 'transfer', 'someone', 'els', 'ur', 'go', '2', 'bahama', 'callfreefon', '08081560665', 'spea
k', 'live', 'oper', 'claim', 'either', 'bahama', 'cruis', 'ofã', '2000', 'cash', '18', 'opt', 'txt', 'x', '0778620011
7', 'someone', 'conact', 'date', 'servic', 'enter', 'phone', 'fanci', 'find', 'call', 'landlin', '09111030116', 'pobox
12n146tf15', 'hi', '07734396839', 'ibh', 'custom', 'loyalti', 'offer', 'new', 'nokia6600', 'mobil', 'ã', '10', 'txtau
ction', 'txt', 'word', 'start', '81151', 'get', '4t', 'sm', 'auction', 'nokia', '7250i', 'get', 'win', 'free', 'aucti
on', 'take', 'part', 'send', 'nokia', '86021', 'hg', 'suite342', '2land', 'row', 'w1jhl', '16', 'call', 'freephon',
'0800', '542', '0578', 'buy', 'space', 'invad', '4', 'chanc', '2', 'win', 'orig', 'arcad', 'game', 'consol', 'press',
'0', 'game', 'arcad', 'std', 'wap', 'charg', 'see', 'o2', 'co', 'uk', 'game', '4', 'term', 'set', 'purchas', 'big',
'brother', 'alert', 'comput', 'select', 'u', '10k', 'cash', '150', 'voucher', 'call', '09064018838', 'ntt', 'po', 'bo
```

```
In [28]: 1 len(spam_words)
```

Out[28]: 11996

```
In [29]: 1 from collections import Counter
2 spam_count =Counter(spam_words)
3 spam_count.most_common(30)
```

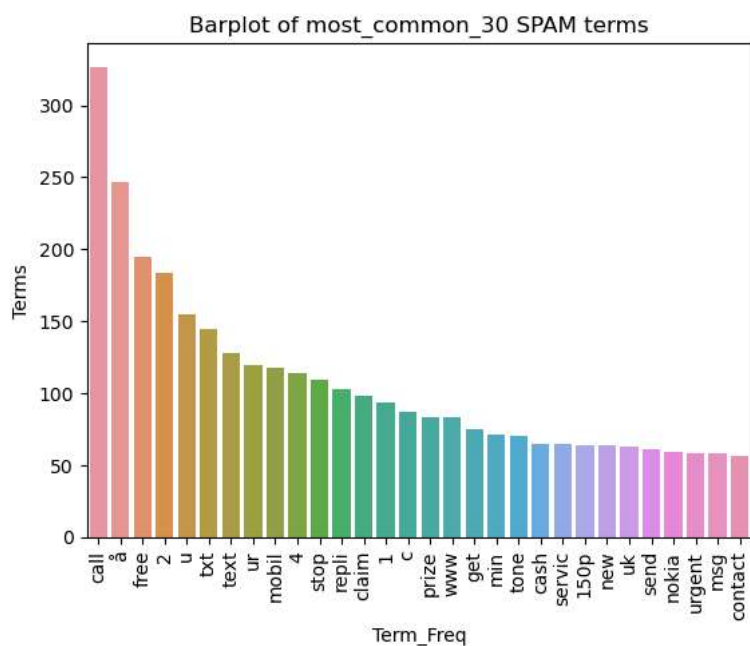
```
Out[29]: [('call', 327),
('ã', 247),
('free', 195),
('2', 184),
('u', 155),
('txt', 145),
('text', 128),
('ur', 119),
('mobil', 118),
('4', 114),
('stop', 109),
('repli', 103),
('claim', 98),
('1', 93),
('c', 87),
('prize', 83),
('www', 83),
('get', 75),
('min', 71),
('tone', 70),
('cash', 65),
('servic', 65),
('150p', 64),
('new', 64),
('uk', 63),
('send', 61),
('nokia', 59),
('urgent', 58),
('msg', 58),
('contact', 56)]
```

```
In [30]: 1 pd.DataFrame(spam_count.most_common(30))
```

Out[30]:

	0	1
0	call	327
1	à	247
2	free	195
3	2	184
4	u	155
5	txt	145
6	text	128
7	ur	119
8	mobil	118
9	4	114
10	stop	109
11	repli	103
12	claim	98
13	1	93
14	c	87
15	prize	83
16	www	83
17	get	75
18	min	71
19	tone	70
20	cash	65
21	servic	65
22	150p	64
23	new	64
24	uk	63
25	send	61
26	nokia	59
27	urgent	58
28	msg	58
29	contact	56

```
In [31]: 1 sns.barplot(x= pd.DataFrame(spam_count.most_common(30))[0],y= pd.DataFrame(spam_count.most_common(30))[1])
2 plt.xticks(rotation = 'vertical')
3 plt.xlabel('Term_Freq')
4 plt.ylabel('Terms')
5 plt.title('Barplot of most_common_30 SPAM terms')
6 plt.show()
```




```
In [32]: 1 # similarly for ham_words
2 ham_words = []
3 for i in df[df.label == 0]['std_text'].str.split():
4     ham_words+=i
5 print(ham_words)
```

'soon', 'come', 'otherwis', 'tomorrow', 'msg', 'r', 'time', 'pa', 'silent', 'say', 'think', 'u', 'right', 'also', 'ma
ke', 'u', 'think', 'least', '4', 'moment', 'gd', 'nt', 'swt', 'drm', 'shesil', 'yeah', 'probabl', 'swing', 'roommat',
'finish', 'girl', 'happi', 'new', 'year', 'melodi', 'ii', 'dun', 'need', 'pick', 'ur', 'gf', 'yay', 'better', 'told',
'5', 'girl', 'either', 'horribl', 'u', 'eat', 'mac', 'eat', 'u', 'forgot', 'abt', 'alreadi', 'rite', 'u', 'take', 'lo
ng', '2', 'repli', 'thk', 'toot', 'b4', 'b', 'prepar', 'wat', 'shall', 'eat', 'say', 'fantast', 'chanc', 'anyth', 'ne
ed', 'bigger', 'life', 'lift', 'lose', '2', 'live', 'think', 'would', 'first', 'person', '2', 'die', 'n', 'v', 'q',
'nw', 'came', 'hme', 'da', 'outsid', 'island', 'head', 'toward', 'hard', 'rock', 'run', 'day', 'class', 'class', 'che
nnai', 'velacheri', 'flippin', 'shit', 'yet', 'k', 'give', 'sec', 'break', 'lt', 'gt', 'cstore', 'much', 'bad', 'avoi
d', 'like', 'yo', 'around', 'got', 'car', 'back', 'annoy', 'goodmorn', 'today', 'late', 'lt', 'gt', 'min', 'point',
'hangin', 'mr', 'right', 'makin', 'u', 'happi', 'come', 'aliv', 'better', 'correct', 'good', 'look', 'fi
gur', 'case', 'guess', 'see', 'campu', 'lodg', 'done', 'come', 'home', 'one', 'last', 'time', 'wont', 'anyth', 'trus
t', 'night', 'worri', 'appt', 'shame', 'miss', 'girl', 'night', 'quiz', 'popcorn', 'hair', 'ok', 'c', 'ya', 'said',
'matter', 'mind', 'say', 'matter', 'al', 'moan', 'n', 'e', 'thin', 'go', 'wrong', 'fault', 'al', 'de', 'argument',
'r', 'fault', 'fed', 'himso', 'bother', 'hav', '2go', 'thanx', 'xx', 'neft', 'transact', 'refer', 'number', 'lt', 'g
t', 'r', 'lt', 'decim', 'gt', 'credit', 'beneficiari', 'account', 'lt', 'gt', 'lt', 'time', 'gt', 'lt', 'gt', 'otherw
is', 'part', 'time', 'job', 'na', 'tuition', 'know', 'call', 'also', 'da', 'feel', 'yesterday', 'night', 'wait', 'ti
l', '2day', 'night', 'dear', 'thank', 'understand', 'tri', 'tell', 'sura', 'whole', 'car', 'appreci', 'last', 'two',
'dad', 'map', 'read', 'semi', 'argument', 'apart', 'thing', 'go', 'ok', 'p', 'need', 'strong', 'arm', 'also', 'maaaa
n', 'miss', 'bday', 'real', 'april', 'guessin', 'gonna', '9', 'ok', 'come', 'ur', 'home', 'half', 'hour', 'yo', 'gam
e', 'almost', 'want', 'go', 'walmart', 'soon', 'yeah', 'probabl', 'sure', 'ilol', 'let', 'u', 'know', 'nerson', 'would

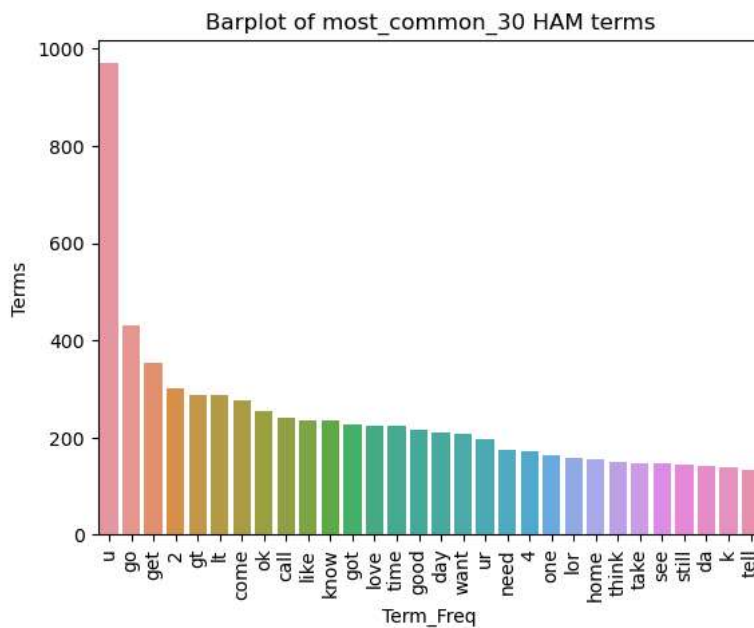
```
In [33]: 1 len(ham_words)
```

Out[33]: 36505

```
In [34]: 1 from collections import Counter
2 ham_count =Counter(ham_words)
3 ham_count.most_common(30)
```

```
Out[34]: [('u', 969),
('go', 432),
('get', 354),
('2', 302),
('gt', 288),
('lt', 287),
('come', 276),
('ok', 255),
('call', 240),
('like', 236),
('know', 236),
('got', 226),
('love', 225),
('time', 223),
('good', 216),
('day', 210),
('want', 209),
('ur', 198),
('need', 174),
('4', 171),
('one', 165),
('lor', 159),
('home', 156),
('think', 150),
('take', 148),
('see', 147),
('still', 145),
('da', 143),
('k', 138),
('tell', 134)]
```

```
In [35]: 1 sns.barplot(x= pd.DataFrame(ham_count.most_common(30))[0],y= pd.DataFrame(ham_count.most_common(30))[1])
2 plt.xticks(rotation = 'vertical')
3 plt.xlabel('Term_Freq')
4 plt.ylabel('Terms')
5 plt.title('Barplot of most_common_30 HAM terms')
6 plt.show()
```



4. Text Vectorization

```
In [36]: 1 from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
2 cv= CountVectorizer(stop_words = 'english')
3 tv= TfidfVectorizer(stop_words = 'english')
4
5 X_cv= cv.fit_transform(df.std_text).toarray()
6 X_tv= tv.fit_transform(df.std_text).toarray()
```

```
In [37]: 1 X_cv.shape, X_tv.shape
```

```
Out[37]: ((5169, 7059), (5169, 7059))
```

- Both the vectorization methods give text vectors of same length

```
In [38]: 1 y=df.label
```

```
In [39]: 1 #Tf-idf vectorization
2 from sklearn.model_selection import cross_val_score, train_test_split
3 xtrain, xtest, ytrain, ytest =train_test_split(X_tv, y, test_size =0.2, random_state =42)
```

5. Model Building

- Which is Better for Spam Detection?
 - If minimizing false positives is critical (i.e., you don't want to risk important emails being marked as spam), precision is more important.
 - This is often the case in professional or business environments where missing an important email can have serious consequences.
 - If minimizing false negatives is critical (i.e., you want to ensure that most spam emails are correctly identified), recall is more important.
 - This is often the case in personal email use where receiving spam emails is more of an annoyance than missing an important email.

```
In [40]: 1 # to check names of scorers for cross-val
2 # import sklearn
3 # sklearn.metrics.get_scorer_names()
```

```
In [41]: 1 from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB
2 from sklearn.model_selection import cross_val_score, train_test_split
3 gnb= GaussianNB()
4 mnb=MultinomialNB()
5 bnb =BernoulliNB()
```

```
In [42]: 1 X={'BoW':X_cv, 'Tf-Idf':X_tv}
2 model={'GaussianNB':gnb, 'MultinomialNB':mnb, 'BernoulliNB':bnb}
```

```
In [43]: 1 scores={}
2         for i in X:
3             for j in model:
4                 scores[i+'_'+j] = np.max(cross_val_score(model[j], X[i], y, cv=5, scoring='precision'))
5
6         # our metric is recall score because, we dont want any false positives (Type 2 error), ie a Spam predicted as Ham
```

```
In [44]: 1 scores
```

```
Out[44]: {'BoW_GaussianNB': 0.5041322314049587,
          'BoW_MultinomialNB': 0.9333333333333333,
          'BoW_BernoulliNB': 1.0,
          'Tf-Idf_GaussianNB': 0.497907949790795,
          'Tf-Idf_MultinomialNB': 1.0,
          'Tf-Idf_BernoulliNB': 1.0}
```

```
In [45]: 1 sorted(scores.items(), key=lambda x: x[1], reverse = True)
```

```
Out[45]: [('BoW_BernoulliNB', 1.0),
          ('Tf-Idf_MultinomialNB', 1.0),
          ('Tf-Idf_BernoulliNB', 1.0),
          ('BoW_MultinomialNB', 0.9333333333333333),
          ('BoW_GaussianNB', 0.5041322314049587),
          ('Tf-Idf_GaussianNB', 0.497907949790795)]
```

- Clearly the winner is BoW with Multinomial Naive Bayes with the best f1_score

```
In [46]: 1 # training the model on 'BoW_BernoulliNB' or 'Tf-Idf_MultinomialNB' or 'Tf-Idf_BernoulliNB' is giving ideal Precisio
2
3         # randomly proceeding with Tf-Idf_MultinomialNB
4         xtrain, xtest, ytrain, ytest = train_test_split(X_tv, y, test_size =0.2, random_state =42)
5         mnb.fit(xtrain, ytrain)
```

```
Out[46]: ▾ MultinomialNB
          MultinomialNB()
```

```
In [47]: 1 from sklearn.metrics import accuracy_score, precision_score, classification_report, confusion_matrix
2         pred_train =mnb.predict(xtrain)
3         pred_test =mnb.predict(xtest)
4         print('training accuracy_score: ', accuracy_score(ytrain, pred_train))
5         print('training precision_score: ', precision_score(ytrain, pred_train))
6         print('training confusion_matrix: \n', confusion_matrix(ytrain, pred_train))
7
8         #predicting testing data
9         print('testing accuracy_score: ', accuracy_score(ytest, pred_test))
10        print('testing precision_score: ', precision_score(ytest, pred_test))
11        print('testing confusion_matrix: \n', confusion_matrix(ytest, pred_test))
```

```
training accuracy_score: 0.9743651753325272
training precision_score: 1.0
training confusion_matrix:
[[3627  0]
 [ 106 402]]
testing accuracy_score: 0.9661508704061895
testing precision_score: 1.0
testing confusion_matrix:
[[889  0]
 [ 35 110]]
```

```
In [48]: 1 # 'Tf-Idf_BernoulliNB'
2         xtrain, xtest, ytrain, ytest =train_test_split(X_tv, y, test_size =0.2, random_state =42)
3         bnb.fit(xtrain, ytrain)
4         from sklearn.metrics import accuracy_score, precision_score, classification_report, confusion_matrix
5         pred_train =bnb.predict(xtrain)
6         pred_test =bnb.predict(xtest)
7         print('training accuracy_score: ', accuracy_score(ytrain, pred_train))
8         print('training precision_score: ', precision_score(ytrain, pred_train))
9         print('training confusion_matrix: \n', confusion_matrix(ytrain, pred_train))
10
11        #predicting testing data
12        print('testing accuracy_score: ', accuracy_score(ytest, pred_test))
13        print('testing precision_score: ', precision_score(ytest, pred_test))
14        print('testing confusion_matrix: \n', confusion_matrix(ytest, pred_test))
```

```
training accuracy_score: 0.9835550181378476
training precision_score: 0.9932735426008968
training confusion_matrix:
[[3624  3]
 [ 65 443]]
testing accuracy_score: 0.971953578336557
testing precision_score: 0.967741935483871
testing confusion_matrix:
[[885  4]
 [ 25 120]]
```

```
In [49]: 1 # 'Bow_BernoulliNB'
2 xtrain, xtest, ytrain, ytest = train_test_split(X_cv, y, test_size = 0.2, random_state = 42)
3 bnb.fit(xtrain, ytrain)
4 from sklearn.metrics import accuracy_score, precision_score, classification_report, confusion_matrix
5 pred_train = bnb.predict(xtrain)
6 pred_test = bnb.predict(xtest)
7 print('training accuracy_score: ', accuracy_score(ytrain, pred_train))
8 print('training precision_score: ', precision_score(ytrain, pred_train))
9 print('training confusion_matrix: \n', confusion_matrix(ytrain, pred_train))
10
11 #predicting testing data
12 print('testing accuracy_score: ', accuracy_score(ytest, pred_test))
13 print('testing precision_score: ', precision_score(ytest, pred_test))
14 print('testing confusion_matrix: \n', confusion_matrix(ytest, pred_test))
```

```
training accuracy_score: 0.9835550181378476
training precision_score: 0.9932735426008968
training confusion_matrix:
[[3624  3]
 [ 65 443]]
testing accuracy_score: 0.971953578336557
testing precision_score: 0.967741935483871
testing confusion_matrix:
[[885  4]
 [ 25 120]]
```

- Since for a Business case '**precision**' is more important than '**accuracy**' for a Spam classifier, we will go ahead with **Tf-Idf-MultinomialNB**.
- Achieved Precision: 1.0
- Accuracy 0.9661508704061895

6. Model Improvement

Trying other ML classifier algorithm

```
In [50]: 1 from sklearn.linear_model import LogisticRegression
2 from sklearn.svm import SVC
3 from sklearn.naive_bayes import MultinomialNB
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.neighbors import KNeighborsClassifier
6 from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, BaggingClassifier, ExtraTreesClassifier, GradientBoostingClassifier
7 from xgboost import XGBClassifier
```

```
In [51]: 1 lr=LogisticRegression()
2 svc =SVC()
3 mnb=MultinomialNB()
4 dt= DecisionTreeClassifier()
5 knn=KNeighborsClassifier()
6 rf= RandomForestClassifier()
7 ada= AdaBoostClassifier()
8 bg=BaggingClassifier()
9 etc= ExtraTreesClassifier()
10 gb= GradientBoostingClassifier()
11 xgb= XGBClassifier()
```

```
In [52]: 1 models ={'LR': lr, 'SVC':svc, 'MultNB':mnb, 'DT':dt, "KNN":knn, "RF":rf, "ADA":ada, "BgC":bg, "ETC":etc, "GB":gb, "XGB":
```

```
In [53]: 1 xtrain, xtest, ytrain, ytest =train_test_split(X_tv, y, test_size =0.2, random_state =42)
2
3 def training_model(model, xtrain, xtest, ytrain, ytest):
4     model.fit(xtrain,ytrain)
5     pred= model.predict(xtest)
6     accuracy =accuracy_score(ytest,pred)
7     precision =precision_score(ytest,pred)
8
9     return accuracy, precision
```

```
In [54]: 1 accuracy=[]
2 precision =[]
3 for i in models:
4     a,p=training_model(models[i],xtrain, xtest, ytrain, ytest)
5     accuracy.append(a)
6     precision.append(p)
```

```
In [55]: 1 final =pd.DataFrame({'Classifier' : models.keys(), 'Accuracy': accuracy, 'Precision': precision})
2 final.sort_values(by='Precision', ascending =False)
```

Out[55]:

	Classifier	Accuracy	Precision
2	MultNB	0.966151	1.000000
4	KNN	0.890716	1.000000
8	ETC	0.979691	0.992063
5	RF	0.977756	0.991935
1	SVC	0.970986	0.991453
9	GB	0.965184	0.990991
0	LR	0.954545	0.980392
10	XGB	0.972921	0.953488
6	ADA	0.974855	0.940741
7	BgC	0.967118	0.930233
3	DT	0.970019	0.895833

Hypertuning the best model

```
In [56]: 1 from sklearn.model_selection import GridSearchCV
```

```
In [57]: 1 from sklearn.pipeline import Pipeline
2 from sklearn.metrics import make_scorer
```

```
In [58]: 1 # hypertuning the TF-IDF -> MNB model
2
3 # Define a pipeline
4 pipeline = Pipeline([
5     ('tfidf', TfidfVectorizer()),
6     ('nb', MultinomialNB())
7 ])
8
9 # Define the parameter grid
10 param_grid = {
11     'tfidf__max_df': [0.5, 0.75, 1.0],
12     'tfidf__min_df': [1, 2, 5],
13     'tfidf__ngram_range': [(1, 1), (1, 2)],
14     'nb__alpha': [0.01, 0.1, 1, 10]
15 }
16
17 # 'tfidf__max_df': Maximum document frequency (proportion of documents in which a term appears).
18 # 'tfidf__min_df': Minimum document frequency (minimum number of documents a term must be in).
19 # 'tfidf__ngram_range': The lower and upper boundary of the range of n-values for different n-grams to be extracted.
20 # 'nb__alpha': Smoothing parameter for the Naive Bayes classifier.
21
22 # Define the scoring dictionary
23 scoring = {
24     'accuracy': make_scorer(accuracy_score),
25     'precision': make_scorer(precision_score)
26 }
27 # Perform GridSearchCV
28 gv = GridSearchCV(pipeline, param_grid, cv=5, n_jobs=-1, verbose=1, scoring = scoring, refit='accuracy')
29
30 # Fit the model
31 xtrain, xtest, ytrain, ytest =train_test_split(df.std_text.astype('str'), y, test_size =0.2, random_state =42)
32 gv.fit(xtrain, ytrain)
33
34 # Print best parameters and best score
35 print("Best parameters found: ", gv.best_params_)
36 print("Best accuracy_score: ", gv.best_score_)
```

Fitting 5 folds for each of 72 candidates, totalling 360 fits

Best parameters found: {'nb__alpha': 0.1, 'tfidf__max_df': 0.5, 'tfidf__min_df': 2, 'tfidf__ngram_range': (1, 1)}

Best accuracy_score: 0.9842805320435308

```
In [59]: 1 # precision, accuracy
2 precision_score(ytest,gv.predict(xtest)), accuracy_score(ytest,gv.predict(xtest))
```

Out[59]: (0.9692307692307692, 0.9777562862669246)

- the accuracy has improved from 0.0.9661508704061895 to 0.9777562862669246 , but precision has dropped from 1.0 to 0.9692307692307692
- this tradeoff is not desirable as precision holds greater significance for business use in a spam classifier
- so we will keep things unchanged and proceed with the **TF-IDF-> MNB** as our final model

```
In [60]: 1 X_tv= tv.fit_transform(df.std_text).toarray()
2 xtrain, xtest, ytrain, ytest =train_test_split(X_tv, y, test_size =0.2, random_state =42)
3 mnb.fit(xtrain, ytrain)
4 pred_train =mnb.predict(xtrain)
5 pred_test =mnb.predict(xtest)
6 print('testing accuracy_score: ', accuracy_score(ytest, pred_test))
7 print('testing precision_score: ', precision_score(ytest, pred_test))
8 print('testing confusion_matrix: \n', confusion_matrix(ytest, pred_test))
```

testing accuracy_score: 0.9661508704061895
testing precision_score: 1.0
testing confusion_matrix:
[[889 0]
[35 110]]

7. Pickle

```
In [61]: 1 import pickle
2 pickle.dump(mnb, open('model.pkl', 'wb'))
3 pickle.dump(tv, open('vectorizer.pkl', 'wb'))
4 pickle.dump(standardize, open('standardize.pkl', 'wb'))
```

```
In [62]: 1 xtest[9].shape
```

Out[62]: (7059,)

```
In [63]: 1 xtest[0].reshape(1, -1).shape
```

Out[63]: (1, 7059)

```
In [64]: 1 mnb.predict(xtest[5].reshape(1, -1)), ytest.iloc[5]
```

Out[64]: (array([1], dtype=int64), 1)

```
In [65]: 1 mnb.predict(xtest[5].reshape(1, -1))[0]
```

Out[65]: 1

```
In [ ]: 1
```

```
In [66]: 1 t= ""We're happy to inform you that you're entitled to a refund for overpayment on your AMEX account. Click on this
```



```
In [67]: 1 standardize(t)
```

Out[67]: 'happi inform entitl refund overpay amex account click link link claim refund'

```
In [68]: 1 tv.transform([standardize(t)]).shape
```

Out[68]: (1, 7059)

```
In [69]: 1 mnb.predict(tv.transform([standardize(t)]))
```

Out[69]: array([0], dtype=int64)

```
In [70]: 1 mnb.predict_proba(tv.transform([standardize(t)]))
```

Out[70]: 0.6703360228456645